

# FHG Case - Predictive Modeling

*Eli (Ilya) Bolotin*

*08/05/2018*

In the first part of this project we dealt with incorrect or missing data by computing and retrieving missing values (where possible). Unfortunately, we were still left with missing values that could not be computed or recovered. To deal with these missing values we will use imputation.

## Load libraries

```
# load VIM
library(VIM)

# load MICE
library(mice)

# load Amelia
library(Amelia)

# load GGPlot
library(ggplot2)

# load forecasting and time series libraries
library(forecast)
library(tseries)
```

## Stage 1: Imputation with MICE and AMELIA

### Pre-imputation analysis

Import our dataset for imputation

```
data <- read.csv("dataset_for_imputation.csv", header=TRUE, sep=",")
```


Find out how many NAs we have as percentage of all rows, for every variable

```
pMiss <- function(x){sum(is.na(x))/length(x)*100}
apply(data,2,pMiss)
```

## Incoming.Examinations	Year	Month
## 8.333333	0.000000	0.000000

Result is 8.33% of observations in the Incoming.Examinations column have NAs. View this in graph form.

```
aggr_plot <- aggr(data, col=c('navyblue','red'), numbers=TRUE, sortVars=TRUE,
  ↪ labels=names(data), ylab=c("Histogram of missing data","Pattern"))
```



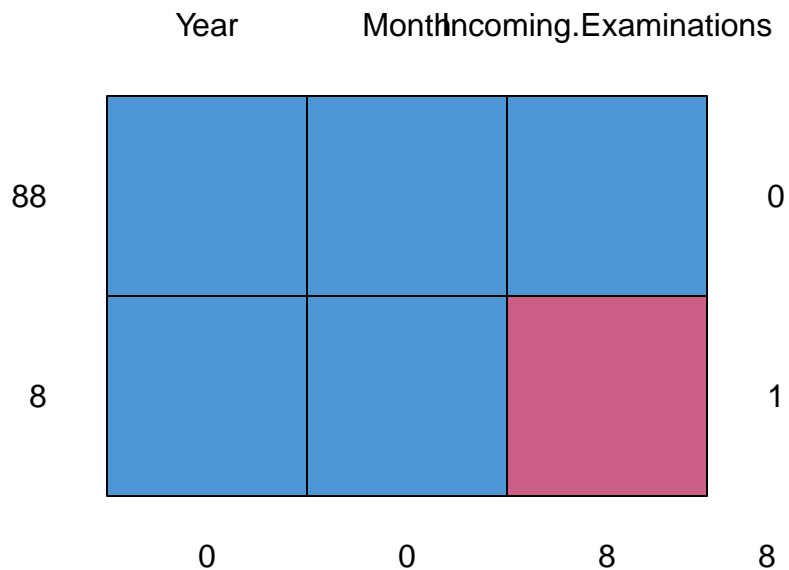
```
##
## Variables sorted by number of missings:
##      Variable      Count
## Incoming.Examinations 0.08333333
##      Year 0.00000000
##      Month 0.00000000
```

To deal with these NAs, we will use imputation (the data has been cleaned previously). In this analysis, 2 forms of imputation are tested. The first is MICE imputation and the second is AMELIA.

## Imputation with MICE

Let's start by getting a better sense of missing data.

```
md.pattern(data)
```



```
##      Year Month Incoming.Examinations
## 88      1      1                1 0
## 8       1      1                0 1
##        0      0                8 8
```

This tells us there are 8 missing values in the Incoming.Examinations column.

Let's impute the data these missing values.

```
# default method is predictive mean matching
imputed_data <- mice(data, m=10, maxit=10, seed=500, print=F)
```

MICE will generate 10 sets of multiple imputations (with 10 iterations per set). We will:

1. Fit each set with a linear model
2. Then select 3 out of 10 sets
3. Pool the fitted datasets into one
4. Then review summary the summary information.

See below:

```
# fit a linear model to the imputed data
mice_fit <- with(imputed_data, lm(Incoming.Examinations ~ Year + Month))

# review summary of linear model for imputation 1
summary(mice_fit$analyses[[1]])
```

```
##
## Call:
```

```
## lm(formula = Incoming.Examinations ~ Year + Month)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -915.6 -458.7 -137.5  380.2 1768.1
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.228e+06  5.155e+04 -23.818  < 2e-16 ***
## Year         6.119e+02  2.566e+01  23.851  < 2e-16 ***
## Month        5.252e+01  1.703e+01   3.084  0.00269 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 576 on 93 degrees of freedom
## Multiple R-squared:  0.8615, Adjusted R-squared:  0.8585
## F-statistic: 289.2 on 2 and 93 DF,  p-value: < 2.2e-16
```

```
# review summary of linear model for imputation 2
summary(mice_fit$analyses[[2]])
```

```
##
## Call:
## lm(formula = Incoming.Examinations ~ Year + Month)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -917.9 -459.9 -138.7  375.9 1767.9
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.226e+06  5.139e+04 -23.86  < 2e-16 ***
## Year         6.109e+02  2.557e+01  23.89  < 2e-16 ***
## Month        5.262e+01  1.697e+01   3.10  0.00256 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 574.1 on 93 degrees of freedom
## Multiple R-squared:  0.8619, Adjusted R-squared:  0.8589
## F-statistic: 290.1 on 2 and 93 DF,  p-value: < 2.2e-16
```

```
# review summary of linear model for imputation 2
summary(mice_fit$analyses[[3]])
```

```
##
## Call:
## lm(formula = Incoming.Examinations ~ Year + Month)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -933.2 -440.8 -148.6  377.9 1760.1
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.222e+06  5.147e+04 -23.739  < 2e-16 ***
```

```
## Year          6.088e+02  2.561e+01  23.771 < 2e-16 ***
## Month         5.592e+01  1.700e+01   3.289  0.00142 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 575 on 93 degrees of freedom
## Multiple R-squared:  0.861, Adjusted R-squared:  0.858
## F-statistic: 287.9 on 2 and 93 DF,  p-value: < 2.2e-16
```

```
# pool the fitted models for all imputed datasets to come up with overall regression
mice_pooled_fit <- pool(mice_fit)
```

```
# adjusted R squared
pool.r.squared(mice_fit, adjusted = TRUE)
```

```
##              est      lo 95      hi 95 fmi
## adj R^2 0.8579779 0.7942441 0.9031312 NaN
```

```
# summarize pooled linear model
summary(mice_pooled_fit)
```

```
##              estimate  std.error  statistic      df    p.value
## (Intercept) -1.224184e+06 51633.75773 -23.708984 90.75259 0.00000000
## Year         6.100202e+02   25.69475  23.741046 90.75261 0.00000000
## Month        5.481914e+01   17.10999   3.203925 90.09275 0.00187097
```

Base on the pooled fit, we can now create a cleaned and completed dataset. Let's do that next.

```
completedDataMice <- complete(imputed_data)
completedDataMice
```

```
## Incoming.Examinations Year Month
## 1          362 2006      1
## 2          436 2006      2
## 3          362 2006      3
## 4          490 2006      4
## 5          508 2006      5
## 6          393 2006      6
## 7          393 2006      7
## 8          596 2006      8
## 9          634 2006      9
## 10         613 2006     10
## 11         545 2006     11
## 12         411 2006     12
## 13         398 2007      1
## 14         311 2007      2
## 15         664 2007      3
## 16         680 2007      4
## 17         442 2007      5
## 18         467 2007      6
## 19         566 2007      7
## 20         806 2007      8
## 21         732 2007      9
## 22         886 2007     10
## 23         776 2007     11
## 24         698 2007     12
```

## 25	875	2008	1
## 26	840	2008	2
## 27	724	2008	3
## 28	1115	2008	4
## 29	997	2008	5
## 30	775	2008	6
## 31	886	2008	7
## 32	1041	2008	8
## 33	1011	2008	9
## 34	775	2008	10
## 35	939	2008	11
## 36	1004	2008	12
## 37	1004	2009	1
## 38	1065	2009	2
## 39	1263	2009	3
## 40	962	2009	4
## 41	1004	2009	5
## 42	1429	2009	6
## 43	1205	2009	7
## 44	890	2009	8
## 45	1320	2009	9
## 46	1276	2009	10
## 47	1757	2009	11
## 48	2043	2009	12
## 49	1491	2010	1
## 50	1595	2010	2
## 51	1578	2010	3
## 52	1604	2010	4
## 53	1758	2010	5
## 54	1595	2010	6
## 55	1457	2010	7
## 56	1607	2010	8
## 57	1808	2010	9
## 58	1866	2010	10
## 59	1934	2010	11
## 60	2294	2010	12
## 61	2294	2011	1
## 62	2334	2011	2
## 63	1973	2011	3
## 64	2262	2011	4
## 65	2259	2011	5
## 66	2217	2011	6
## 67	2739	2011	7
## 68	2772	2011	8
## 69	3383	2011	9
## 70	2869	2011	10
## 71	2239	2011	11
## 72	2789	2011	12
## 73	2789	2012	1
## 74	3455	2012	2
## 75	2940	2012	3
## 76	2968	2012	4
## 77	3466	2012	5
## 78	3037	2012	6


```
## 79          3946 2012    7
## 80          3459 2012    8
## 81          3446 2012    9
## 82          3258 2012   10
## 83          4729 2012   11
## 84          3694 2012   12
## 85          4610 2013    1
## 86          4841 2013    2
## 87          5172 2013    3
## 88          4351 2013    4
## 89          5187 2013    5
## 90          4710 2013    6
## 91          5010 2013    7
## 92          4978 2013    8
## 93          5008 2013    9
## 94          6094 2013   10
## 95          4874 2013   11
## 96          5933 2013   12
```

Create CSV with cleaned, completed data.

```
write.csv(completedDataMice, "cleaned_dataset.csv", row.names=F)
```

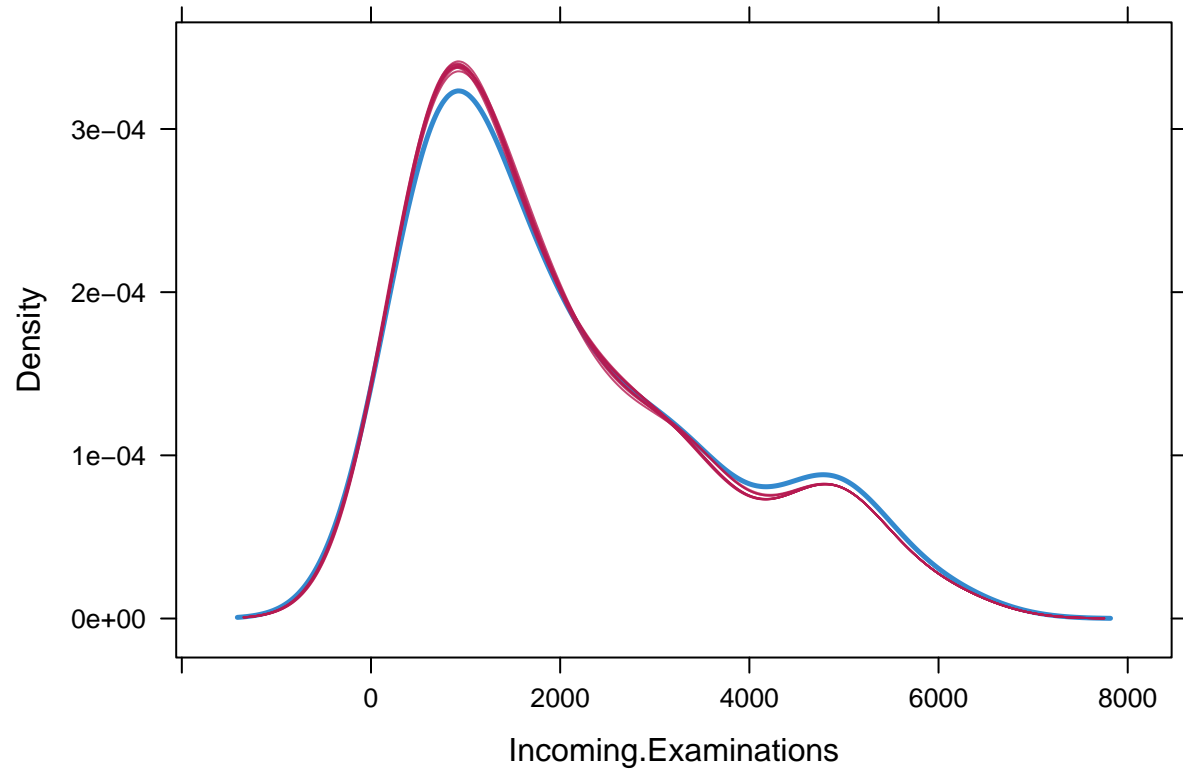
Compare the (cleaned) imputed data to the original data.

```
# blue points are observed, red are imputed. The overlap tells us that the imputed values
↪ are plausible
xyplot(imputed_data, Incoming.Examinations ~ Year, pch=18,cex=2)
```

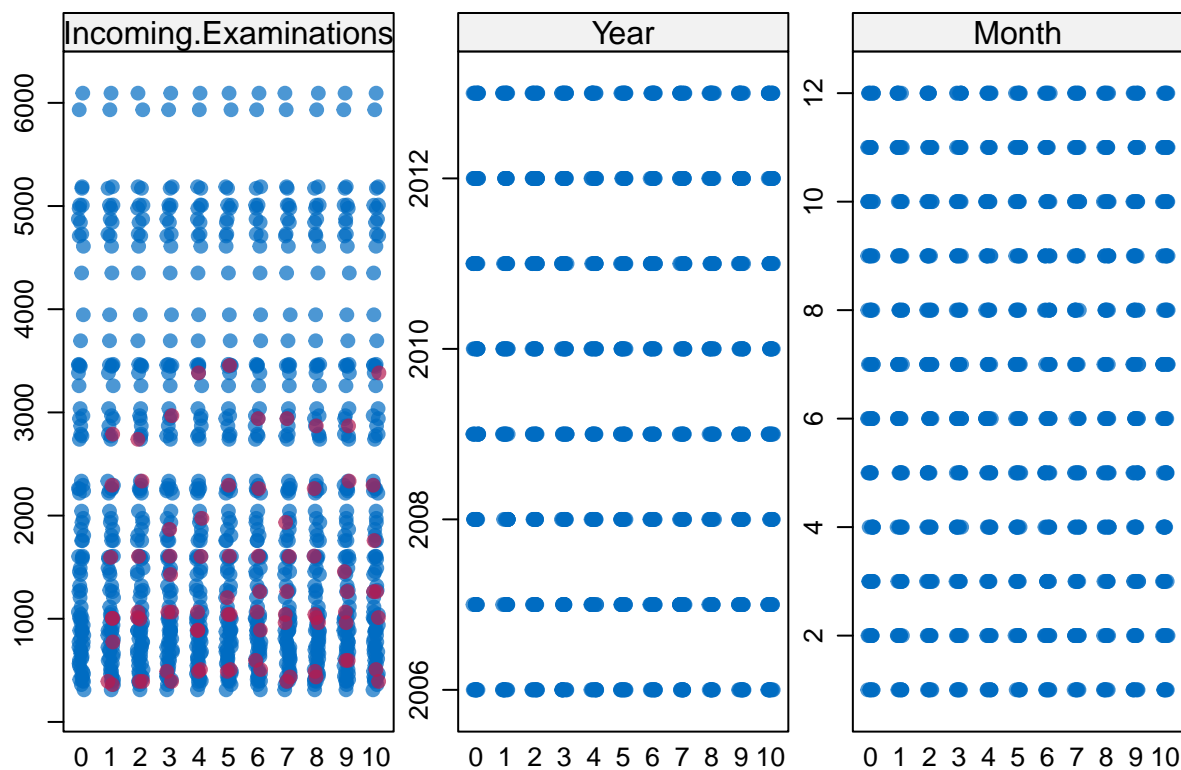


```
# Density plot of the imputed dataset (96 observations)  
densityplot(imputed_data, n=96)
```





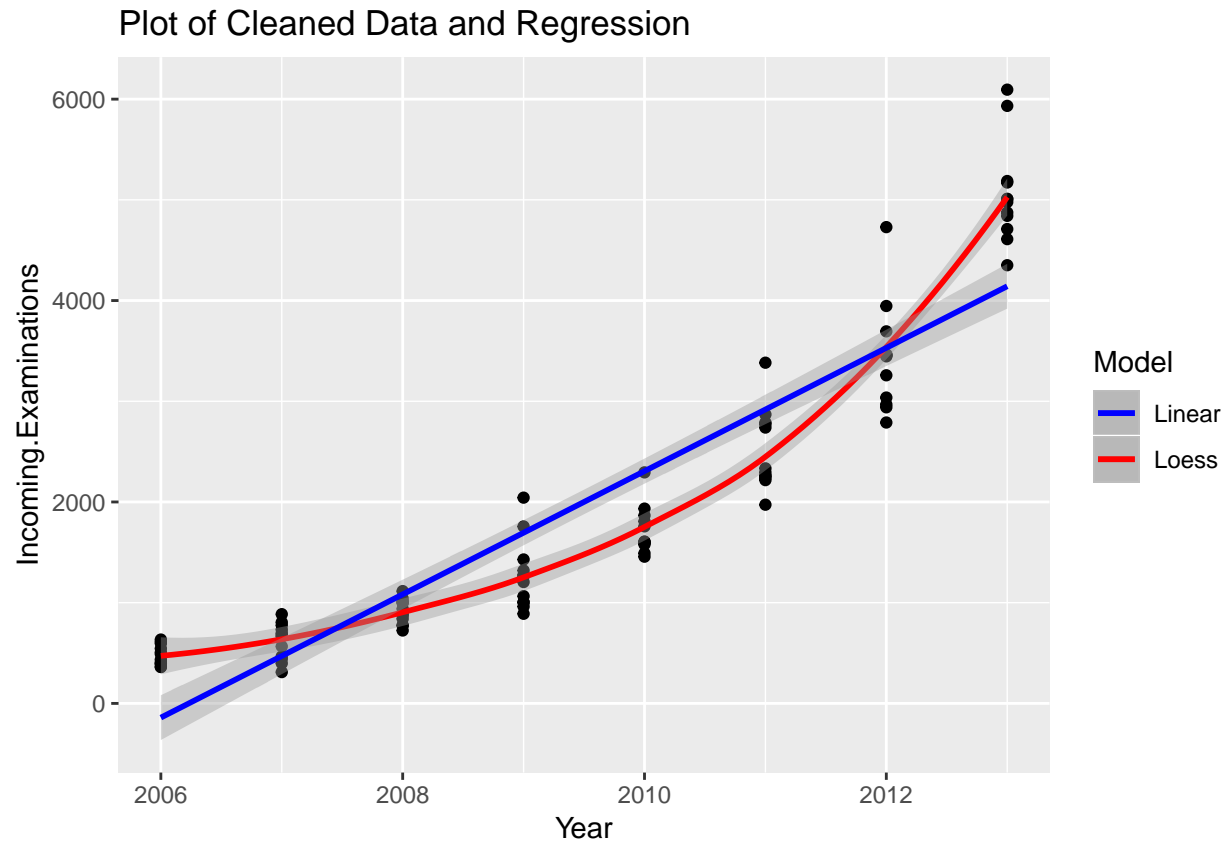
```
# Distributions of the imputed values for each variable by imputed dataset  
stripplot(imputed_data, pch = 20, cex = 1.2)
```



As shown in the charts above, our imputed values are both plausible and follow a similar distribution as the other values in set.

Next, we plot the linear and loess models against our imputed dataset (set 1) to check which fits better.

```
ggplot(completedDataMice, aes(x = Year, y = Incoming.Examinations)) +
  geom_point() +
  geom_smooth(method="loess", aes(colour="Loess")) +
  geom_smooth(method="lm", aes(colour="Linear")) +
  ggtitle("Plot of Cleaned Data and Regression") +
  scale_colour_manual(name="Model", values=c("blue", "red"))
```



As you can see, the observations do not follow a linear model. Nonlinear regression would be better suited to fit this data.

## Imputation with AMELIA

In this section, we will test another method of value imputation to compare it with MICE imputation above.

Let's import our dataset:

```
data <- read.csv("dataset_for_imputation.csv", header=TRUE, sep=",")
```

Run AMELIA imputation:

```
amelia_imp <- amelia(data, m=10, parallel = "multicore", ts="Year", p2s=0)
```

Convert list to dataframe and round values.

```
amelia_imp <- amelia_imp$imputations
amelia_ds <- do.call(rbind.data.frame, amelia_imp)
```

*# round values*

```
amelia_ds <- round(amelia_ds[,], 0)
```

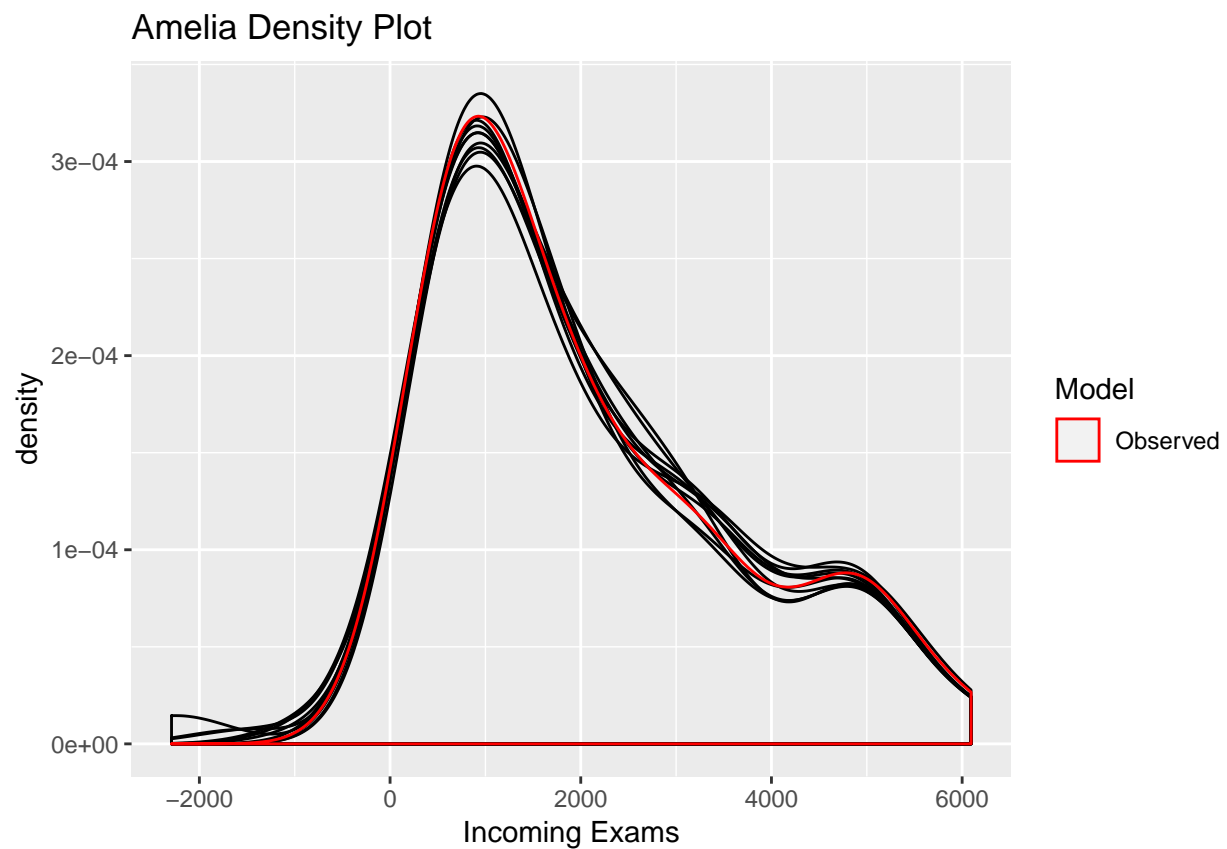
Evaluate the distributions of Amelia imputations

```
ggplot(amelia_imp[[1]]) +
  geom_density(aes(x=amelia_imp[[1]]$Incoming.Examinations)) +
  geom_density(aes(x=amelia_imp[[2]]$Incoming.Examinations)) +
```

```

geom_density(aes(x=amelia_imp[[3]]$Incoming.Examinations)) +
geom_density(aes(x=amelia_imp[[4]]$Incoming.Examinations)) +
geom_density(aes(x=amelia_imp[[5]]$Incoming.Examinations)) +
geom_density(aes(x=amelia_imp[[6]]$Incoming.Examinations)) +
geom_density(aes(x=amelia_imp[[7]]$Incoming.Examinations)) +
geom_density(aes(x=amelia_imp[[8]]$Incoming.Examinations)) +
geom_density(aes(x=amelia_imp[[9]]$Incoming.Examinations)) +
geom_density(aes(x=amelia_imp[[10]]$Incoming.Examinations)) +
geom_density(aes(x=data$Incoming.Examinations, col="Observed")) +
scale_colour_manual(name="Model", values=c("red")) +
labs(title="Amelia Density Plot", x="Incoming Exams")

```



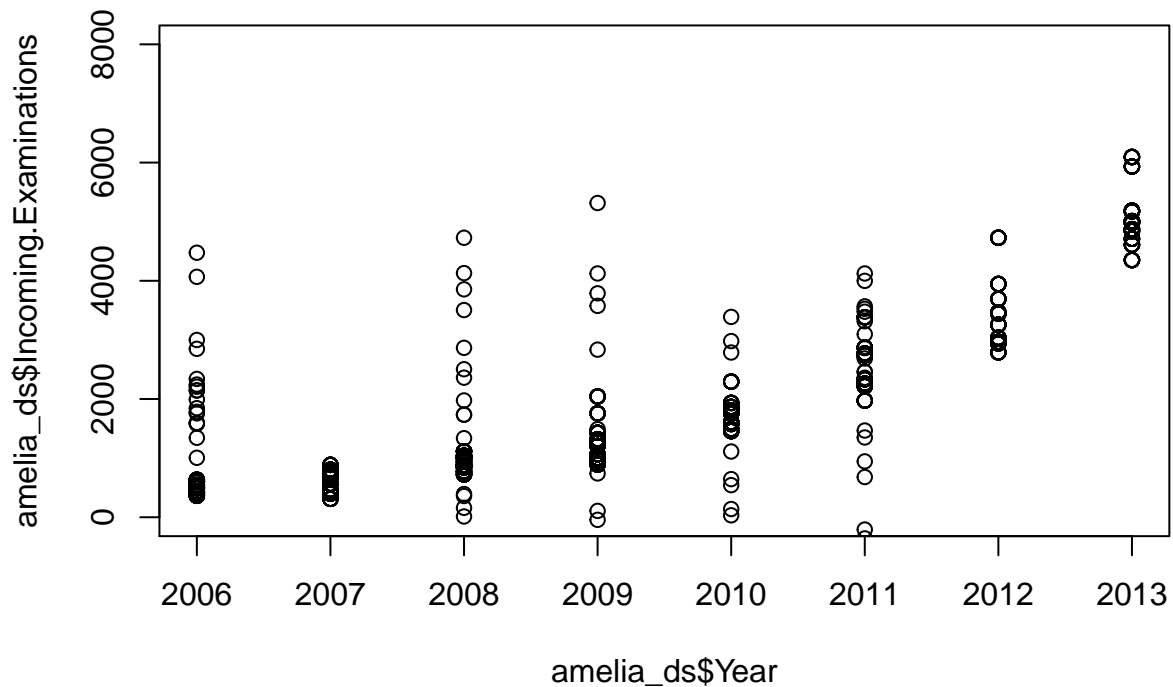
Clearly, every imputation follows a similar distribution.

Plot our dataset with the (AMELIA) imputed observations.

```

plot(amelia_ds$Incoming.Examinations ~ amelia_ds$Year, ylim=c(0,8000))

```



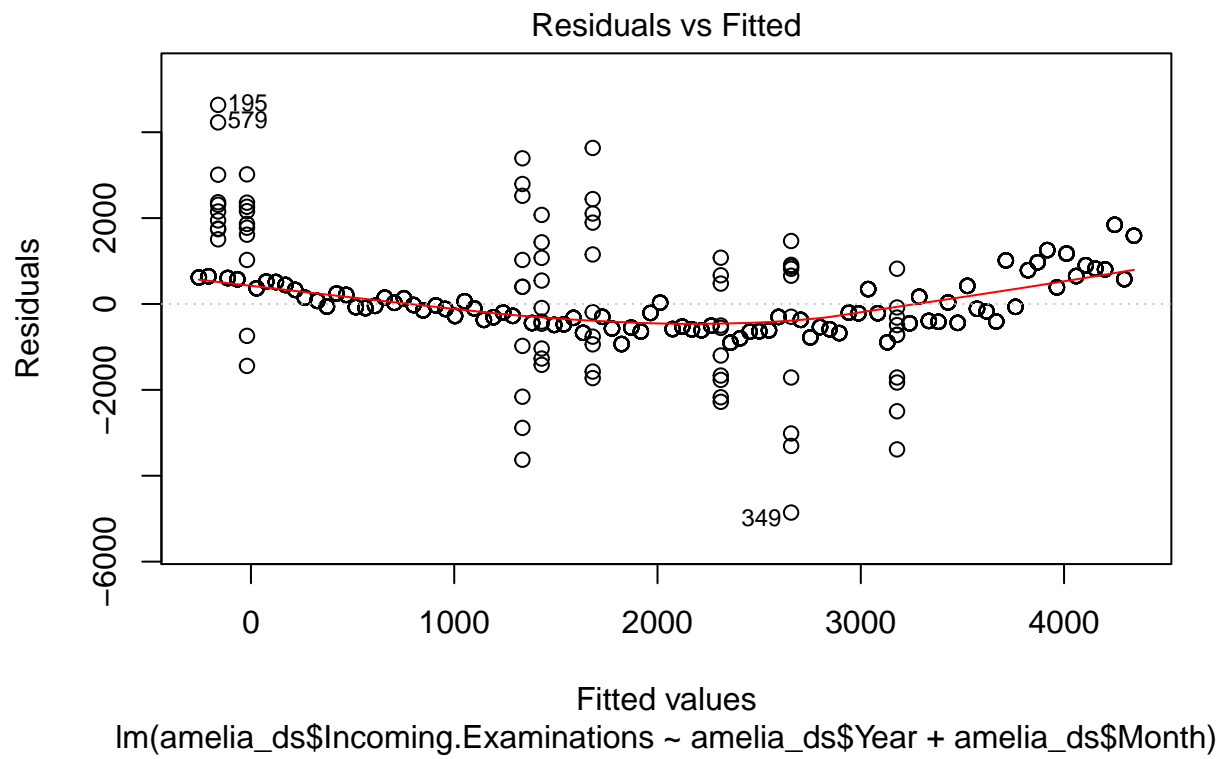
Next, we would like to fit a regression on Amelia imputations, then analyze this regression, and finally, plot regression.

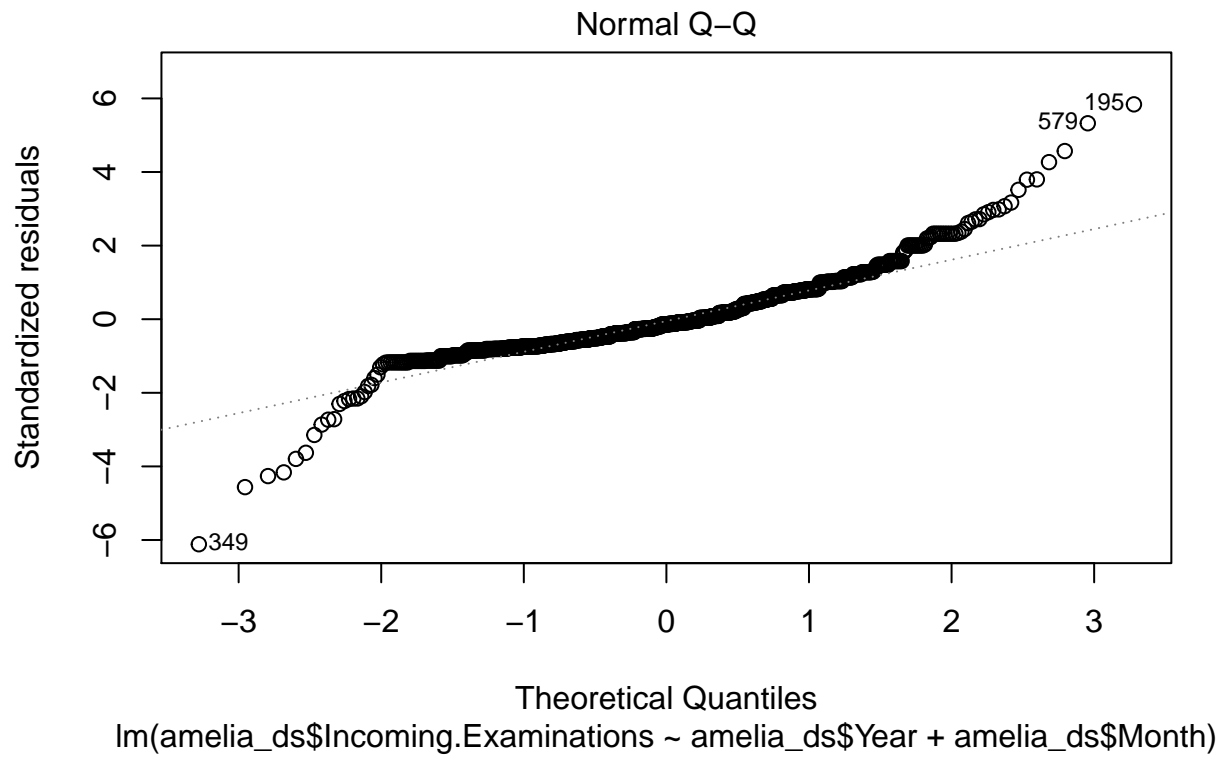
```
# run a linear fit on the first imputed Amelia dataset
amelia_fit <- with(amelia_ds, lm(amelia_ds$Incoming.Examinations ~ amelia_ds$Year +
  ↪ amelia_ds$Month)) #
summary(amelia_fit)
```

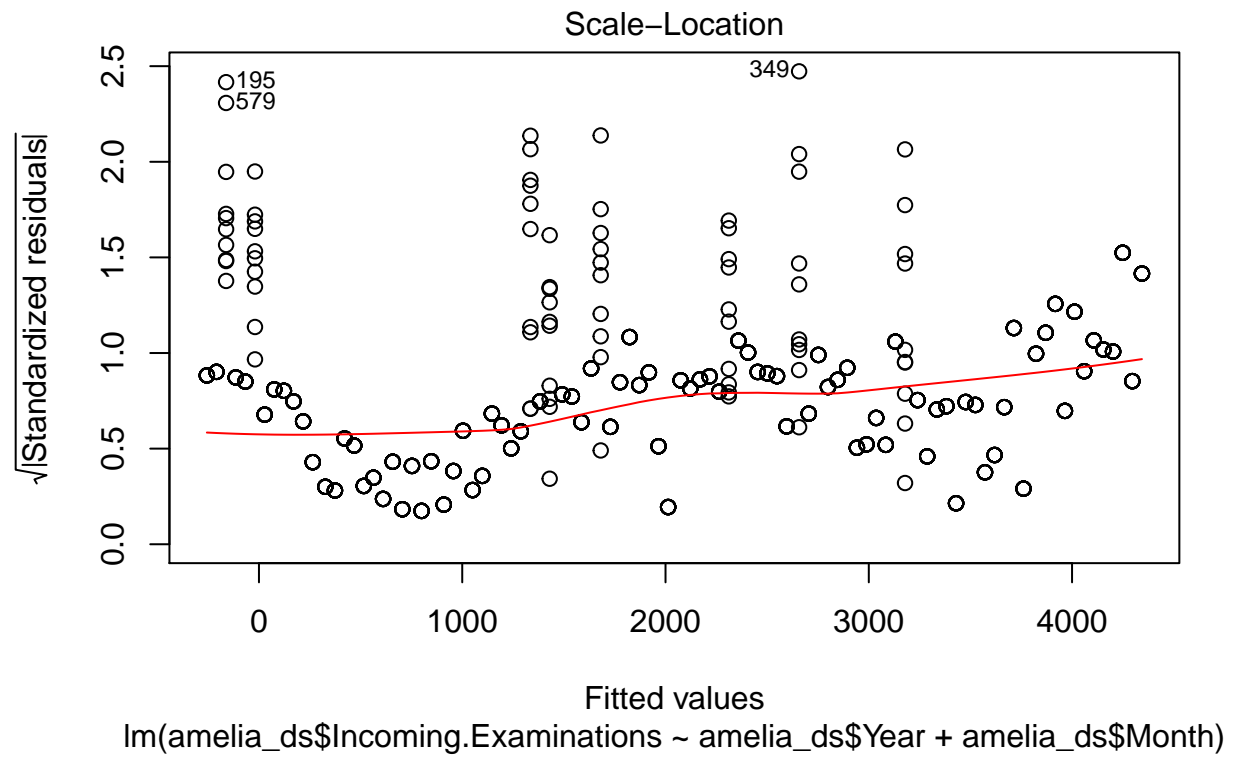
```
##
## Call:
## lm(formula = amelia_ds$Incoming.Examinations ~ amelia_ds$Year +
##     amelia_ds$Month)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4857.2  -487.8  -112.6   405.8  4636.6
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -1.169e+06  2.253e+04  -51.899  < 2e-16 ***
## amelia_ds$Year   5.827e+02  1.121e+01   51.976  < 2e-16 ***
## amelia_ds$Month   4.738e+01  7.441e+00   6.367  2.99e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 795.9 on 957 degrees of freedom
## Multiple R-squared:  0.7413, Adjusted R-squared:  0.7407
```

```
## F-statistic: 1371 on 2 and 957 DF, p-value: < 2.2e-16
```

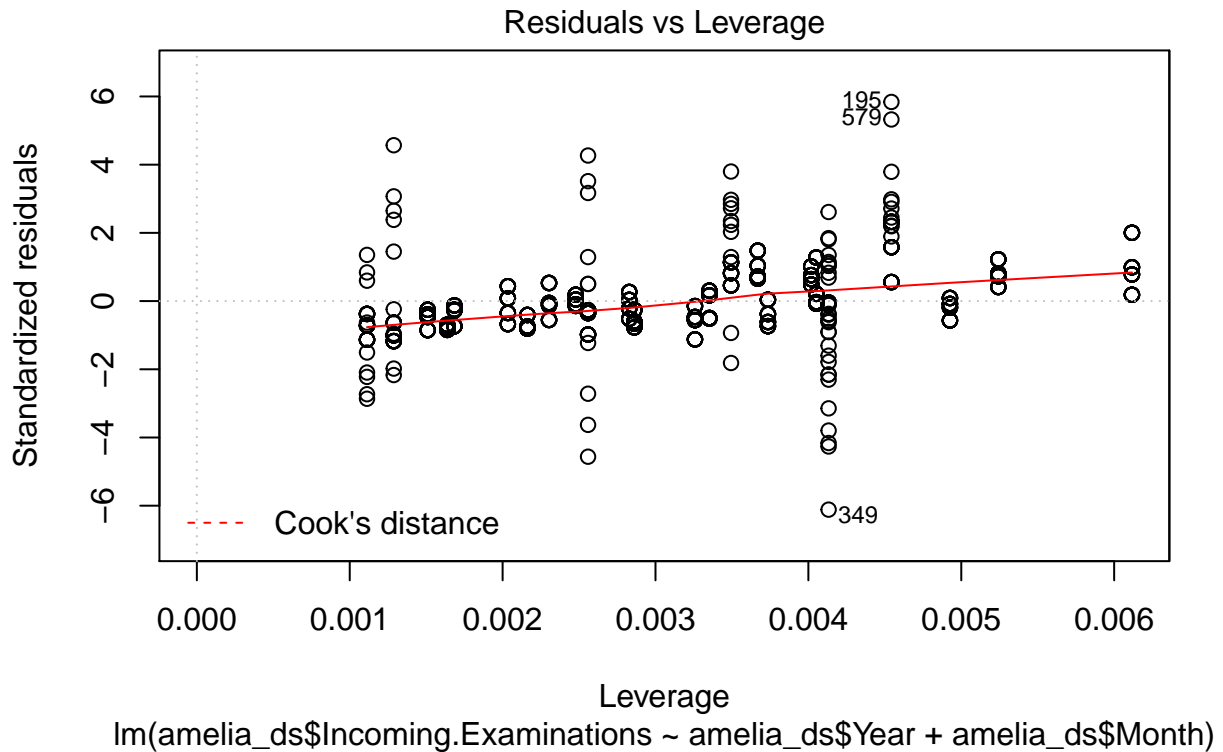
```
plot(amelia_fit)
```











## Stage 2: Forecasting


Now that we have dealt with missing values by using imputation, our next goal is to forecast demand of incoming examinations. In this stage, we will use ARIMA and Holt's exponential smoothing to generate predictive models.

### Forecasting with ARIMA (with MICE imputed dataset)

Step 1: review time series and ensure that the time series is stationary.


```
# create time series
ts = ts(completedDataMice$Incoming.Examinations, start=c(2006, 1), end=c(2013,12),
        frequency=12)

# plot the time series
plot(ts, xlab="Time", ylab="Incoming Examinations")
```




```
# decompose the time series to understand the prevalence of three components: trend,  
↪ seasonality, and error/irregularity  
components.ts = decompose(ts)  
  
# we see that the data is not stationary  
plot(components.ts)
```

## Decomposition of additive time series




```
# we can test this another way by using seasonal decomposition in loess smoothing  
plot(stl(log(ts), s.window="period"))
```



*# autocorrelation function shows data is not stationary. Correlates set of observations  
 ↳ at current time to set of observations at k periods earlier.*  
 acf(ts)

## Series ts



```
# run ADF test to confirm our hypothesis that data is not stationary  
adf.test(ts)
```

```
##  
## Augmented Dickey-Fuller Test  
##  
## data: ts  
## Dickey-Fuller = -0.31208, Lag order = 4, p-value = 0.9884  
## alternative hypothesis: stationary
```


```
# estimate the number of differences required to make a given time series stationary  
ndiffs(ts)
```

```
## [1] 1
```

```
# differentiate the data to make it stationary. We differentiate it twice to remove what  
↪ appears to be a quadratic trend.
```

```
diff_ts <- diff(ts, differences=2)
```

```
# take a look at the differentiated data  
plot(diff_ts)
```




```
# run ADF to confirm data is now stationary
adf.test(diff_ts)
```


```
##
## Augmented Dickey-Fuller Test
##
## data: diff_ts
## Dickey-Fuller = -8.4792, Lag order = 4, p-value = 0.01
## alternative hypothesis: stationary
```

Step 2: identify reasonable model or models (possible values of p and q)

```
# review the ACF chart of the differenced time series. This plots the AC at various lags
↳ for the stationary/differenced time series.
Acf(diff_ts)
```



```
# correlation between an observation at current period and an observation at k periods  
↳ earlier with observations between removed.  
Pacf(diff_ts)
```



Step 3: fit the model

*# (p, d, q). d = 2 for nonseasonal differences applied. The lag at which the PACF cuts  
 ↳ off is the indicated number of AR terms (p). The lag at which the ACF cuts off is the  
 ↳ indicated number of MA terms (q).*

```
fit <- Arima(ts, order=c(7,2,6))
```

*# Automated forecasting using an ARIMA model*

```
fitauto <- auto.arima(ts)
```

*# lower  $S^2$ , meaning points are closer to the line*

```
summary(fit)
```

```
## Series: ts
## ARIMA(7,2,6)
##
## Coefficients:
##      ar1      ar2      ar3      ar4      ar5      ar6      ar7      ma1
##    -0.5466  0.2136  0.606   0.900  -0.1045  -0.2040  0.1025  -1.3558
## s.e.   0.1277  0.1717  0.176   0.111   0.1452   0.1465  0.1300   0.1674
##      ma2      ma3      ma4      ma5      ma6
##    -0.1624  -0.0394  0.1334  1.3699  -0.9419
## s.e.   0.3626   0.3858  0.3599  0.3970   0.1980
##
## sigma^2 estimated as 86696:  log likelihood=-668.53
## AIC=1365.06  AICc=1370.37  BIC=1400.66
##
```



```
## Training set error measures:
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 34.12339 270.4629 195.3399 -0.8952236 12.34182 0.291775
##           ACF1
## Training set -0.03209728
```

```
summary(fitauto)
```

```
## Series: ts
## ARIMA(1,1,1) with drift
##
## Coefficients:
##           ar1      ma1      drift
##          -0.2949 -0.4919  54.9360
## s.e.    0.1428   0.1282  12.8896
##
## sigma^2 estimated as 103157:  log likelihood=-681.93
## AIC=1371.87   AICc=1372.31   BIC=1382.09
##
## Training set error measures:
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.6192313 314.4183 222.5433 -7.53879 15.81654 0.3324082
##           ACF1
## Training set 9.147416e-05
```

```
# correlation of custom ARIMA fitted values to actual values
cor(fitted(fit),ts)^2
```

```
## [1] 0.9691572
```


```
# correlation of custom ARIMA fitted values to actual values
cor(fitted(fitauto),ts)^2
```

```
## [1] 0.9602115
```

Step 4: evaluate the model's fit, accuracy and residuals

```
# checking distribution & Ljung-Box for ARIMA
# chart a Q-Q plot to test if the data is normally distributed
qqnorm(fit$residuals)
qqline(fit$residuals) # add line
```

## Normal Q-Q Plot



*# the results of this test are not significant, suggesting that the autocorrelations  
↪ don't differ from zero.*

```
Box.test(fit$residuals, type="Ljung-Box")
```


```
##  
## Box-Ljung test  
##  
## data: fit$residuals  
## X-squared = 0.10203, df = 1, p-value = 0.7494
```

*# checking distribution & Ljung-Box for auto ARIMA*

```
qqnorm(fitauto$residuals)
```

```
qqline(fitauto$residuals)
```

## Normal Q-Q Plot



```
Box.test(fitauto$residuals, type="Ljung-Box") # lowest chi-square suggesting better fit,
↳ and highest p-value suggesting low AC and nonsignificance. But this is due to mean
↳ error.
```

```
##
## Box-Ljung test
##
## data: fitauto$residuals
## X-squared = 8.2865e-07, df = 1, p-value = 0.9993
```

```
# checking accuracy
accuracy(fit) # check accuracy for ARIMA fit.
```

```
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 34.12339 270.4629 195.3399 -0.8952236 12.34182 0.291775
##           ACF1
## Training set -0.03209728
```

```
accuracy(fitauto) # check accuracy for ARIMA auto fit
```

```
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.6192313 314.4183 222.5433 -7.53879 15.81654 0.3324082
##           ACF1
## Training set 9.147416e-05
```


```
# checking residuals
summary(residuals(fit))
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.    Max.
## -684.2587 -119.7795  -0.0792   34.1234  142.3442  925.1034
```


```
summary(residuals(fitauto))
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.    Max.
## -883.1740 -193.6513  -42.5549  -0.6192  100.6924 1161.3995
```


```
boxplot(residuals(fit))
```




```
boxplot(residuals(fitauto))
```




```
plot(residuals(fit))
```



```
plot(residuals(fitauto))
```



```
# compare ARIMA models to time series  
plot(ts)  
lines(fitted(fit), col="green", lwd="4")  
lines(fitted(fitauto), col="red", lwd="2")
```



Step 5: make forecasts and show predictions


```
# create forecast for custom ARIMA model
forecast_arima <- forecast(fit, h=12)

# create forecast for auto ARIMA model
forecast_arimaauto <- forecast(fitauto, h=12)

# plot custom ARIMA forecast
plot(forecast_arima)
```




### Forecasts from ARIMA(7,2,6)



```
# plot auto generated forecast  
plot(forecast_arimaauto)
```

## Forecasts from ARIMA(1,1,1) with drift



```
# show forecasted incoming exams for next 12 months using custom ARIMA model
forecast_arima
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2014	5648.928	5264.640	6033.216	5061.211	6236.646
## Feb 2014	6204.625	5818.351	6590.898	5613.870	6795.379
## Mar 2014	5966.686	5564.426	6368.946	5351.482	6581.890
## Apr 2014	6534.957	6132.134	6937.779	5918.892	7151.021
## May 2014	6459.474	6051.289	6867.659	5835.209	7083.738
## Jun 2014	6681.212	6267.628	7094.796	6048.689	7313.735
## Jul 2014	6797.003	6380.862	7213.145	6160.570	7433.437
## Aug 2014	7143.707	6713.717	7573.698	6486.093	7801.322
## Sep 2014	7105.961	6670.395	7541.526	6439.821	7772.100
## Oct 2014	7352.476	6902.050	7802.903	6663.609	8041.344
## Nov 2014	7588.955	7130.822	8047.087	6888.302	8289.607
## Dec 2014	7750.958	7268.747	8233.169	7013.480	8488.436

```
# show forecasted incoming exams for next 12 months using the auto ARIMA model
forecast_arimaauto
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2014	5499.183	5087.574	5910.793	4869.681	6128.686
## Feb 2014	5698.272	5277.412	6119.132	5054.622	6341.922
## Mar 2014	5710.692	5251.658	6169.725	5008.660	6412.723
## Apr 2014	5778.167	5293.639	6262.696	5037.145	6519.190
## May 2014	5829.405	5318.063	6340.747	5047.375	6611.435
## Jun 2014	5885.432	5349.355	6421.509	5065.573	6705.290

```
## Jul 2014      5940.046 5380.119 6499.974 5083.711 6796.382
## Aug 2014      5995.077 5412.333 6577.821 5103.847 6886.308
## Sep 2014      6049.985 5445.268 6654.703 5125.150 6974.821
## Oct 2014      6104.930 5479.015 6730.844 5147.675 7062.184
## Nov 2014      6159.863 5513.444 6806.282 5171.251 7148.475
## Dec 2014      6214.800 5548.508 6881.092 5195.794 7233.805
```

## Forecasting with Holt's (with MICE imputed dataset)

```
# generate time series
ts = ts(completedDataMice$Incoming.Examinations, start=c(2006, 1), end=c(2013,12),
↪ frequency=12)

# Holt's approach but with multiplicative error and multiplicative trend.
fitholts_m <- ets(ts, model="MMN")

# Holt's model (additive)
fitholts_a <- ets(ts, model="AAN")

# check Holt's multiplicative accuracy.
accuracy(fitholts_m)

##                ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  9.777434 269.1309 198.3447 -0.629297 12.30327 0.2962633
##                ACF1
## Training set -0.1297505

# check Holt's additive model accuracy
accuracy(fitholts_a)


##                ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 67.84433 298.2579 208.7351 0.2865199 13.20034 0.3117831
##                ACF1
## Training set -0.009329275

# create Holts multiplicative forecast
forecast_m <- forecast(fitholts_m, 12)

# create Holts additive forecast for comparison
forecast_a <- forecast(fitholts_a, 12)


# plot Holts multiplicative forecast
plot(forecast_m, main="(Holt's Multiplicative) Forecast for Incoming Examinations",
↪ ylab="Inc. Exams", xlab="Months", flty=2)
```

## (Holt's Multiplicative) Forecast for Incoming Examinations



```
# plot Holts forecast
plot(forecast_a, main="(Holt's Additive) Forecast for Incoming Examinations", ylab="Inc.
↳ Exams", xlab="Months", flty=2)
```

## (Holt's Additive) Forecast for Incoming Examinations



```
# show forecasted incoming exams for next 12 months by Holts_m model
forecast_m
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2014	5998.765	4745.877	7213.437	4177.846	7824.393
## Feb 2014	6175.419	4911.081	7463.346	4206.967	8074.614
## Mar 2014	6357.274	5044.117	7649.880	4332.045	8342.059
## Apr 2014	6544.485	5183.895	7866.999	4465.023	8563.324
## May 2014	6737.209	5345.225	8114.385	4649.733	8831.467
## Jun 2014	6935.608	5530.809	8330.830	4741.420	9118.043
## Jul 2014	7139.850	5695.271	8609.541	4856.907	9396.467
## Aug 2014	7350.106	5820.520	8856.302	5071.857	9625.589
## Sep 2014	7566.554	5997.112	9143.919	5169.957	9927.934
## Oct 2014	7789.376	6258.686	9387.505	5425.262	10167.111
## Nov 2014	8018.760	6370.831	9655.999	5532.910	10554.243
## Dec 2014	8254.898	6530.365	9922.261	5645.952	10847.938


```
# show forecasted incoming exams for next 12 months by Holts_a model
forecast_a
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2014	5841.241	5450.788	6231.695	5244.094	6438.389
## Feb 2014	6001.130	5610.291	6391.968	5403.394	6598.866
## Mar 2014	6161.018	5769.316	6552.721	5561.961	6760.076
## Apr 2014	6320.907	5927.673	6714.141	5719.507	6922.306
## May 2014	6480.795	6085.181	6876.410	5875.755	7085.835
## Jun 2014	6640.684	6241.666	7039.701	6030.439	7250.928


```
## Jul 2014      6800.572 6396.969 7204.175 6183.315 7417.829
## Aug 2014      6960.460 6550.946 7369.975 6334.162 7586.759
## Sep 2014      7120.349 6703.473 7537.225 6482.792 7757.906
## Oct 2014      7280.237 6854.448 7706.026 6629.049 7931.425
## Nov 2014      7440.126 7003.796 7876.455 6772.817 8107.434
## Dec 2014      7600.014 7151.463 8048.565 6914.014 8286.014
```

Check residuals & summaries of Holt's multiplicative and additive models.

```
qqnorm(fitholts_m$residuals)
qqline(fitholts_m$residuals)
```



```
qqnorm(fitholts_a$residuals)
qqline(fitholts_a$residuals)
```



```
summary(fitholts_m$residuals)
```

```
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
## -0.42203 -0.07770 -0.01506  0.01810  0.11162  0.41333
```

```
summary(fitholts_a$residuals)
```

```
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
## -662.17  -88.55   29.49   67.84  189.29  940.47
```